

assignment-3-slcak

October 19, 2023

ASSIGNMENT_3

```
[ ]: import pandas as pd
```

2. Load the dataset.

```
[ ]: # Load a CSV dataset
data = pd.read_csv(r"C:/Users/sonudr/Downloads/penguins_size.csv")
```

```
[ ]: data.head()
```

```
[ ]: species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1             18.7             181.0
1  Adelie  Torgersen         39.5             17.4             186.0
2  Adelie  Torgersen         40.3             18.0             195.0
3  Adelie  Torgersen          NaN             NaN              NaN
4  Adelie  Torgersen         36.7             19.3             193.0

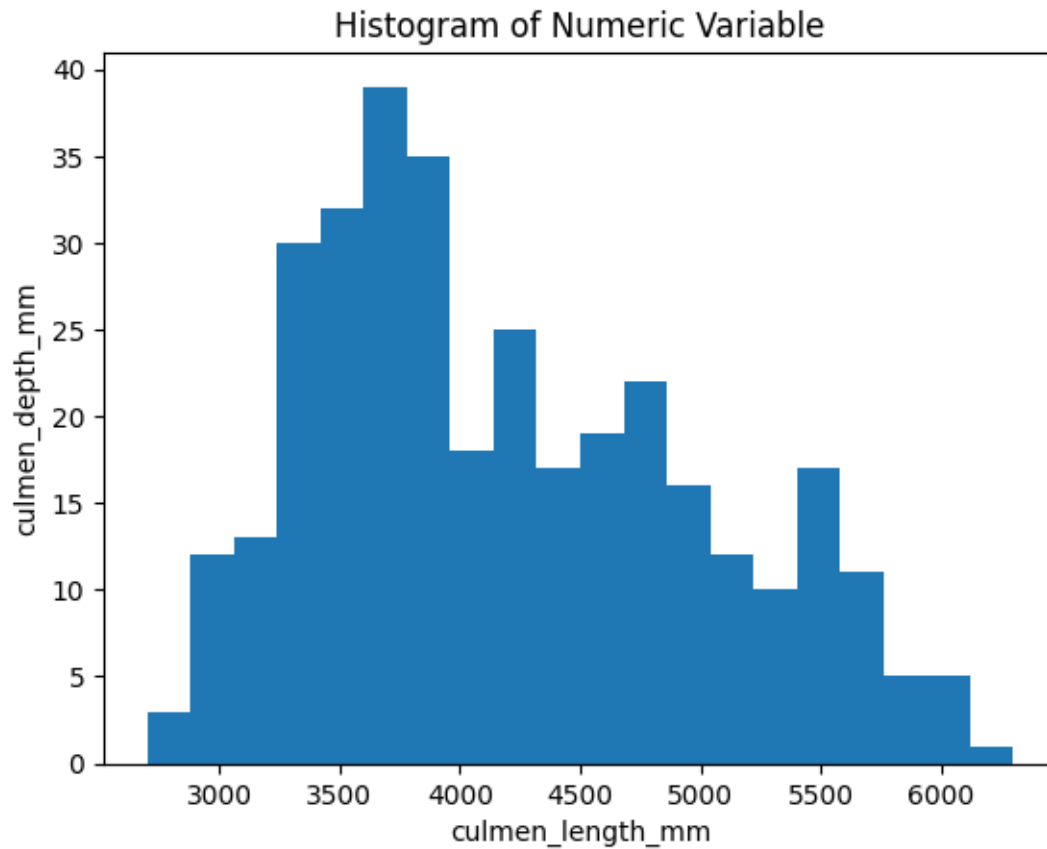
      body_mass_g      sex
0         3750.0    MALE
1         3800.0  FEMALE
2         3250.0  FEMALE
3            NaN      NaN
4         3450.0  FEMALE
```

3. Perform the Below Visualizations.

Univariate Analysis:

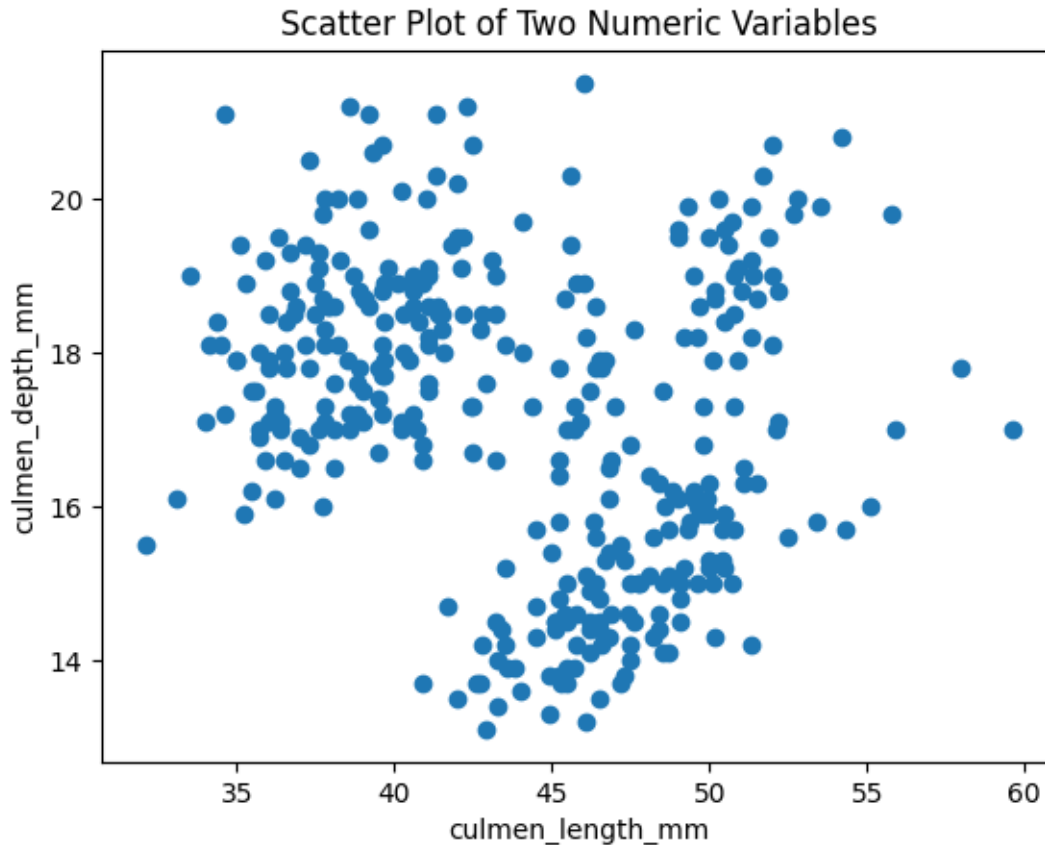
```
[ ]: import matplotlib.pyplot as plt

# Plot a histogram for a numeric variable
plt.hist(data['body_mass_g'], bins=20)
plt.xlabel('culmen_length_mm')
plt.ylabel('culmen_depth_mm')
plt.title('Histogram of Numeric Variable')
plt.show()
```



Bivariate Analysis:

```
[ ]: # Scatter plot for two numeric variables
plt.scatter(data['culmen_length_mm'], data['culmen_depth_mm'])
plt.xlabel('culmen_length_mm')
plt.ylabel('culmen_depth_mm')
plt.title('Scatter Plot of Two Numeric Variables')
plt.show()
```

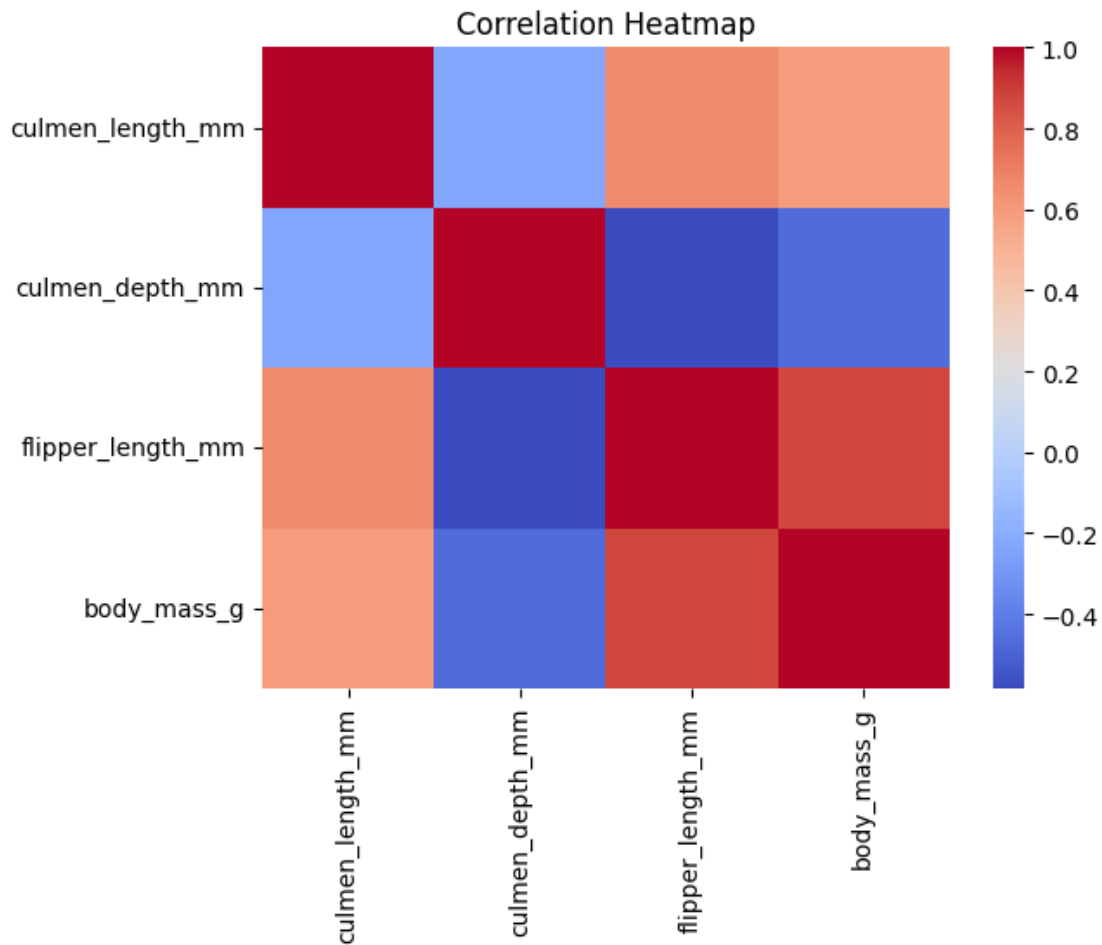


Multivariate Analysis:

```
[ ]: # Compute the correlation matrix
import seaborn as sns
correlation_matrix = data.corr()

# Create a heatmap to visualize correlations
sns.heatmap(correlation_matrix, annot=False, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```

```
C:\Users\sonudr\AppData\Local\Temp\ipykernel_6608\2299829469.py:3:
FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
    correlation_matrix = data.corr()
```



```
[ ]: import pandas as pd

# Assuming 'data' is your DataFrame
numeric_data = data.select_dtypes(include=['number']) # Select numeric columns

# Calculate descriptive statistics
descriptive_stats = numeric_data.describe()

# Print the results
print(descriptive_stats)
```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
count	342.000000	342.000000	342.000000	342.000000
mean	43.921930	17.151170	200.915205	4201.754386
std	5.459584	1.974793	14.061714	801.954536
min	32.100000	13.100000	172.000000	2700.000000
25%	39.225000	15.600000	190.000000	3550.000000
50%	44.450000	17.300000	197.000000	4050.000000

75%	48.500000	18.700000	213.000000	4750.000000
max	59.600000	21.500000	231.000000	6300.000000

```
[ ]: data.head()
```

```
[ ]: species      island  culmen_length_mm  culmen_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen         39.1           18.7           181.0
1  Adelie  Torgersen         39.5           17.4           186.0
2  Adelie  Torgersen         40.3           18.0           195.0
3  Adelie  Torgersen          NaN           NaN           NaN
4  Adelie  Torgersen         36.7           19.3           193.0

    body_mass_g  sex
0         3750.0  MALE
1         3800.0  FEMALE
2         3250.0  FEMALE
3             NaN   NaN
4         3450.0  FEMALE
```

4. Perform descriptive statistics on the dataset.

```
[ ]: # Calculate mean for a specific column
mean_value = data['body_mass_g'].mean()
print("Mean:", mean_value)

# Calculate median for a specific column
median_value = data['body_mass_g'].median()
print("Median:", median_value)

# Calculate standard deviation for a specific column
std_deviation = data['body_mass_g'].std()
print("Standard Deviation:", std_deviation)
```

Mean: 4201.754385964912

Median: 4050.0

Standard Deviation: 801.9545356980956

5. Handle the Missing values.

```
[ ]: missing_values = data.isnull().sum()
print(missing_values)
```

```
species      0
island        0
culmen_length_mm  2
culmen_depth_mm  2
flipper_length_mm  2
body_mass_g    2
sex           10
```

dtype: int64

Remove Rows with Missing Values:

```
[ ]: data_cleaned = data.dropna()
      print(data_cleaned)
```

	species	island	culmen_length_mm	culmen_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7	181.0	
1	Adelie	Torgersen	39.5	17.4	186.0	
2	Adelie	Torgersen	40.3	18.0	195.0	
4	Adelie	Torgersen	36.7	19.3	193.0	
5	Adelie	Torgersen	39.3	20.6	190.0	
..	
338	Gentoo	Biscoe	47.2	13.7	214.0	
340	Gentoo	Biscoe	46.8	14.3	215.0	
341	Gentoo	Biscoe	50.4	15.7	222.0	
342	Gentoo	Biscoe	45.2	14.8	212.0	
343	Gentoo	Biscoe	49.9	16.1	213.0	

	body_mass_g	sex
0	3750.0	MALE
1	3800.0	FEMALE
2	3250.0	FEMALE
4	3450.0	FEMALE
5	3650.0	MALE
..
338	4925.0	FEMALE
340	4850.0	FEMALE
341	5750.0	MALE
342	5200.0	FEMALE
343	5400.0	MALE

[334 rows x 7 columns]

Impute Missing Values - Numeric Variables:

```
[ ]: mean_value = data['flipper_length_mm'].mean()
      data['flipper_length_mm'].fillna(mean_value, inplace=True)
```

6. Find the outliers and replace them outliers

```
[ ]: import pandas as pd
      import numpy as np

      # Create the DataFrame as previously mentioned
      num_observations = 4
      num_features = 4
      data = np.random.rand(num_observations, num_features)
```

```

df = pd.DataFrame(data, columns=['culmen_length_mm', 'culmen_depth_mm', 'flipper_length_mm', 'body_mass_g'])

# Function to replace outliers with the median
def replace_outliers_with_median(column):
    Q1 = column.quantile(0.25)
    Q3 = column.quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    column = column.apply(lambda x: x if lower_bound <= x <= upper_bound else column.median())
    return column

# Replace outliers in 'culmen_length_mm' column
df['culmen_length_mm'] = replace_outliers_with_median(df['culmen_length_mm'])

# Display the DataFrame with replaced outliers
print(df)

```

	culmen_length_mm	culmen_depth_mm	flipper_length_mm	body_mass_g
0	0.871495	0.919863	0.604437	0.672170
1	0.803383	0.562585	0.196712	0.764539
2	0.428299	0.844605	0.119821	0.679133
3	0.264061	0.740250	0.210042	0.213816

7. Check the correlation of independent variables with the target

```

[ ]: data = {
    'culmen_length_mm': np.random.rand(100),
    'culmen_depth_mm': np.random.rand(100),
    'body_mass_g': np.random.rand(100),
    'flipper_length_mm': np.random.rand(100)
}

df = pd.DataFrame(data)

# Calculate the correlation matrix
correlation_matrix = df.corr()

# Extract the correlation of independent variables with the flipper_length_mm
correlation_with_flipper_length_mm = correlation_matrix['flipper_length_mm'].
    drop('flipper_length_mm') # Remove the flipper_length_mm's self-correlation

# Display the correlation with the flipper_length_mm
print(correlation_with_flipper_length_mm)

```

culmen_length_mm	0.107236
------------------	----------

```
culmen_depth_mm    -0.031346
body_mass_g        0.000361
Name: flipper_length_mm, dtype: float64
```

8. Check for Categorical columns and perform encoding.

```
[ ]: # Perform one-hot encoding
df_encoded = pd.get_dummies(df, columns=['body_mass_g', 'flipper_length_mm'])

# Display the encoded DataFrame
print(df_encoded)
```

	culmen_length_mm	culmen_depth_mm	body_mass_g_0.009974733101252742	\
0	0.619497	0.788435		0
1	0.061157	0.955190		0
2	0.030322	0.836502		0
3	0.863425	0.005652		0
4	0.786718	0.576331		0
..	
95	0.386928	0.799389		0
96	0.471303	0.282308		0
97	0.217129	0.114600		0
98	0.254672	0.009530		0
99	0.724172	0.066864		0

	body_mass_g_0.03324159140960847	body_mass_g_0.04981787625357004	\
0			1
1			0
2			0
3			0
4			0
..	
95			0
96			0
97			0
98			0
99			0

	body_mass_g_0.05325829753231048	body_mass_g_0.0634358119503885	\
0			0
1			0
2			0
3			0
4			0
..	
95			0
96			0
97			0

98	0	0
99	0	0

	body_mass_g_0.06348806559949194	body_mass_g_0.06835124837536233	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
..	
95	0	0	
96	0	0	
97	0	0	
98	0	0	
99	1	0	

	body_mass_g_0.0809441235883206	...	flipper_length_mm_0.9235590499580286	\
0	0	...	0	
1	0	...	0	
2	0	...	0	
3	0	...	0	
4	0	...	0	
..	
95	0	...	0	
96	0	...	0	
97	0	...	0	
98	0	...	0	
99	0	...	0	

	flipper_length_mm_0.9373638237566263	\
0	0	
1	0	
2	0	
3	0	
4	0	
..	...	
95	0	
96	0	
97	0	
98	0	
99	0	

	flipper_length_mm_0.9538769914143631	\
0	0	
1	0	
2	0	
3	0	
4	0	

```

..
95
96
97
98
99

flipper_length_mm_0.9621805244880077 \
0
1
2
3
4
..
95
96
97
98
99

flipper_length_mm_0.9687033782107564 \
0
1
2
3
4
..
95
96
97
98
99

flipper_length_mm_0.9792281438203934 \
0
1
2
3
4
..
95
96
97
98
99

flipper_length_mm_0.9829535717675453 \
0

```

1	0
2	0
3	0
4	0
..	...
95	0
96	0
97	0
98	0
99	0

	flipper_length_mm_0.9904843879367103 \
0	0
1	0
2	0
3	0
4	0
..	...
95	0
96	0
97	0
98	0
99	0

	flipper_length_mm_0.9906799993503721	flipper_length_mm_0.9923362375458921
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
..
95	0	0
96	0	0
97	0	0
98	0	1
99	0	0

[100 rows x 202 columns]

```
[ ]: from sklearn.preprocessing import LabelEncoder

# Initialize the LabelEncoder
label_encoder = LabelEncoder()

# Apply label encoding to the 'Category' column
df['Category_encoded'] = label_encoder.fit_transform(df['body_mass_g'])
```

```
# Display the DataFrame with the encoded 'Category' column
print(df)
```

	culmen_length_mm	culmen_depth_mm	body_mass_g	flipper_length_mm \
0	0.619497	0.788435	0.049818	0.982954
1	0.061157	0.955190	0.905512	0.650588
2	0.030322	0.836502	0.247442	0.374358
3	0.863425	0.005652	0.475928	0.833851
4	0.786718	0.576331	0.914381	0.065310
..
95	0.386928	0.799389	0.595300	0.366489
96	0.471303	0.282308	0.820338	0.311509
97	0.217129	0.114600	0.779467	0.870575
98	0.254672	0.009530	0.712111	0.992336
99	0.724172	0.066864	0.063488	0.279538

	Category_encoded
0	2
1	91
2	30
3	51
4	93
..	...
95	61
96	80
97	75
98	66
99	5

[100 rows x 5 columns]

9. Split the data into dependent and independent variables.

```
[ ]: # Assuming 'culmen_depth_mm' is your target variable (dependent variable)
# and the rest of the columns are your independent variables (features)

# Independent variables (features)
X = df.drop('culmen_depth_mm', axis=1)

# Dependent variable (target)
y = df['culmen_length_mm']

# Display the independent and dependent variables
print("Independent Variables (Features):\n", X)
print("\nDependent Variable (Target):\n", y)
```

Independent Variables (Features):

culmen_length_mm	body_mass_g	flipper_length_mm	Category_encoded
------------------	-------------	-------------------	------------------

0	0.619497	0.049818	0.982954	2
1	0.061157	0.905512	0.650588	91
2	0.030322	0.247442	0.374358	30
3	0.863425	0.475928	0.833851	51
4	0.786718	0.914381	0.065310	93
..
95	0.386928	0.595300	0.366489	61
96	0.471303	0.820338	0.311509	80
97	0.217129	0.779467	0.870575	75
98	0.254672	0.712111	0.992336	66
99	0.724172	0.063488	0.279538	5

[100 rows x 4 columns]

Dependent Variable (Target):

0	0.619497
1	0.061157
2	0.030322
3	0.863425
4	0.786718
...	
95	0.386928
96	0.471303
97	0.217129
98	0.254672
99	0.724172

Name: culmen_length_mm, Length: 100, dtype: float64

10. Scaling the data

```
[ ]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler, StandardScaler

# Assuming you have already split the data into independent variables (X) and
↳ dependent variable (y)

# Min-Max Scaling
min_max_scaler = MinMaxScaler()
X_min_max_scaled = min_max_scaler.fit_transform(X)

# Standardization (Z-score scaling)
standard_scaler = StandardScaler()
X_standard_scaled = standard_scaler.fit_transform(X)

# Display the scaled data
print("Min-Max Scaled Data:\n", pd.DataFrame(X_min_max_scaled, columns=X.
↳ columns))
```

```
print("\nStandardized (Z-score Scaled) Data:\n", pd.
↳DataFrame(X_standard_scaled, columns=X.columns))
```

Min-Max Scaled Data:

	culmen_length_mm	body_mass_g	flipper_length_mm	Category_encoded
0	0.618533	0.040415	0.990401	0.020202
1	0.058162	0.908398	0.650366	0.919192
2	0.027215	0.240877	0.367761	0.303030
3	0.863349	0.472644	0.837857	0.515152
4	0.786363	0.917395	0.051582	0.939394
..
95	0.385119	0.593731	0.359711	0.616162
96	0.469801	0.822001	0.303463	0.808081
97	0.214701	0.780543	0.875430	0.757576
98	0.252381	0.712220	1.000000	0.666667
99	0.723589	0.054282	0.270754	0.050505

[100 rows x 4 columns]

Standardized (Z-score Scaled) Data:

	culmen_length_mm	body_mass_g	flipper_length_mm	Category_encoded
0	0.439466	-1.457500	1.603638	-1.645531
1	-1.515245	1.414238	0.464832	1.437674
2	-1.623196	-0.794267	-0.481637	-0.675534
3	1.293445	-0.027462	1.092756	0.051964
4	1.024898	1.444003	-1.540550	1.506960
..
95	-0.374740	0.373156	-0.508597	0.398392
96	-0.079348	1.128391	-0.696979	1.056604
97	-0.969198	0.991226	1.218589	0.883390
98	-0.837760	0.765178	1.635787	0.571605
99	0.805928	-1.411623	-0.806524	-1.541602

[100 rows x 4 columns]

11. Split the data into training and testing

```
[ ]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

12. check the training and testing data shape.

```
[ ]: # Display the shapes of the resulting sets
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
```

```
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (80, 4)
X_test shape: (20, 4)
y_train shape: (80,)
y_test shape: (20,)
```

```
[ ]:
```